

# Assignment-3

Shivansh Garg (2024MCS2450), Rohit Patidar (2024JCS2042)

April 2025

## Introduction

In this assignment, we add the following capabilities to xv6:

- Memory Printer: This prints the number of pages allocated to a user process.
- Page Swapping: Pages are periodically moved to the disk using an adaptive page replacement policy.

## Modified Disk Layout

The current disk layout of xv6 is:

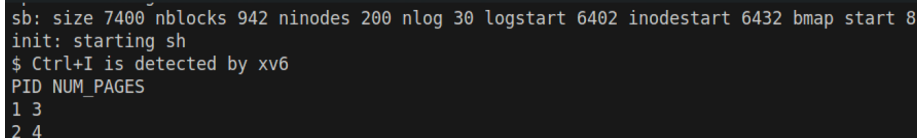
```
[ boot block | sb block | log | inode blocks | free bitmap  
  | data blocks ]
```

To transfer memory pages to disk, a new partition must be introduced. Add a **swap blocks partition** between the **sb** block and the **log** of the disk layout (refer to `mkfs.c`).

Each swap block should be divided into **swap slots**, where each slot represents eight consecutive disk blocks (since a page is 4096 bytes and one disk block is 512 bytes).

Initialize the swap slot array at boot time. You can allocate 800 slots for this purpose and accordingly adjust `FSSIZE` in `param.h`.

When saving a memory page to disk, store the contents into 8 disk blocks directly. Unlike file system writes, **swap block writes bypass the log layer** (refer to `fs.c`) as they handle volatile memory.



```
sb: size 7400 nblocks 942 ninodes 200 nlog 30 logstart 6402 inodestart 6432 bmap start 8  
init: starting sh  
$ Ctrl+I is detected by xv6  
PID NUM_PAGES  
1 3  
2 4
```

Figure 1:

## Memory Organization

The function `kalloc()` is used to allocate memory pages in `xv6`. To restrict physical memory to 4MB, use:

(placed in `memlayout.h`).

To handle memory overflows, implement page swapping logic in a new file like `pageswap.c`.

## Page Replacement

To identify a process and its memory page for swapping (victim process and page):

- Select the process with the highest number of user pages in RAM.
- Add a new field `rss` in `struct proc` to track resident set size.
- If two processes have equal `rss`, pick the one with the lower PID.

When a process exits, clean up its associated swap slots.

To select a victim page:

- Choose a page with the `PTE_P` flag set, and `PTE_A` flag unset.
- The `PTE_P` indicates the page is present in memory.
- The `PTE_A` (accessed flag) is set by the hardware if the page is accessed.

## Swapping Mechanism (Algorithm)

The adaptive swapping mechanism works as follows:

1. **Monitor Free Pages:** Continuously check the number of free pages using the kernel's free space list (e.g., via `get_free_pages()` in kernel space).
2. **Trigger Condition:** If the number of free pages falls below  $Th$ , trigger the following steps:

Print "Current Threshold =  $Th$ , Swapping  $N_{pg}$  pages"
3. **Swap-Out Operation:** Swap out  $N_{pg}$  pages to the disk. Typically, these are selected using LRU or similar page replacement policies.
4. **Update Threshold ( $Th$ ):**

$$Th = \left\lfloor Th \cdot \left(1 - \frac{\beta}{100}\right) \right\rfloor$$

5. Update Swapping Pages ( $N_{pg}$ ):

$$N_{pg} = \left\lfloor \min(LIMIT, N_{pg} \cdot \left(1 + \frac{\alpha}{100}\right)) \right\rfloor$$

```
rohit@rohit:~/Desktop/Sharing_asg3/Sharing_asg3/xv6-onecfile$ make qemu
qemu-system-i386 -serial mon:stdio -drive file=fs.img,index=1,media=disk,format=raw -dr
xv6...
free_pages at the time of kinit2 : 733Swap: Initialized 800 swap slots
cpu0: starting 0
sb: size 7400 nblocks 942 ninodes 200 nlog 30 logstart 6402 inodestart 6432 bmap start 8
init: starting sh
$ memtest
Current Threshold = 100, Swapping 4 pages
Current Threshold = 90, Swapping 5 pages
Current Threshold = 81, Swapping 6 pages
Current Threshold = 73, Swapping 7 pages
Current Threshold = 66, Swapping 8 pages
Current Threshold = 60, Swapping 10 pages
Current Threshold = 54, Swapping 12 pages
Current Threshold = 49, Swapping 15 pages
Current Threshold = 45, Swapping 18 pages
Current Threshold = 41, Swapping 22 pages
Current Threshold = 37, Swapping 27 pages
Current Threshold = 34, Swapping 33 pages
Memtest Passed
```

Figure 2:

## Role of $\alpha$ and $\beta$ in System Efficiency

### $\alpha$ (Alpha) – Growth Factor for $N_{pg}$

- Higher  $\alpha$  leads to a rapid increase in the number of pages swapped out.
- Helps quickly free up large memory chunks under high memory pressure.
- Risk: May cause excessive swapping, increasing I/O and latency.

### $\beta$ (Beta) – Decay Factor for $Th$

- Higher  $\beta$  leads to a faster reduction in the threshold  $Th$ .
- Reduces frequency of future swaps by making the trigger less sensitive.
- Risk: Might delay swap triggers, risking memory exhaustion.

### Combined Effect

Let:

- $\alpha \in [0, 100]$  (growth rate for number of pages to be swapped out)
- $\beta \in [0, 100]$  (decay rate for threshold)
- $Th = 100$  (initial threshold for free pages)
- $N_{pg} = 2$  (initial number of pages to swap out)
- $LIMIT = 100$  (maximum value for  $N_{pg}$ )

We monitor the number of free pages in memory. If the available free pages fall below the current threshold ( $Th$ ), we trigger the swapping mechanism. The goal is to manage memory efficiently without affecting system responsiveness.

- **Low  $\alpha$ , Low  $\beta$ :** Conservative behavior; small swaps, frequent triggers.
- **High  $\alpha$ , Low  $\beta$ :** Aggressive memory freeing with conservative triggering.
- **Low  $\alpha$ , High  $\beta$ :** Conservative swapping but progressively rarer.
- **High  $\alpha$ , High  $\beta$ :** Aggressive and adaptive—best for dynamic systems.

#### Example Configuration:

- Initial values:  $\alpha = 25$ ,  $\beta = 10$
- $N_{pg}$  increases by 25% after each swap
- $Th$  decreases by 10% after each swap

This configuration makes the system increasingly aggressive in freeing memory, while reducing the chance of repeated swaps.

## Makefile Integration

The values of  $\alpha$  and  $\beta$  can be set at compile time using the Makefile:

```
CFLAGS += -DALPHA=25 -DBETA=10
```

In the code, use:

```
int alpha = ALPHA;  
int beta = BETA;
```

This approach enables easy tuning of the parameters without changing the source code.